

Build Models and Analyse FD001

In [84]:

```
1 import os
2 os.environ["KERAS_BACKEND"] = "torch"
3 os.environ["PYTORCH_ENABLE_MPS_FALLBACK"] = "1"
4 os.environ["PYTORCH_MPS_HIGH_WATERMARK_RATIO"] = "0.0" # optional: reduce MPS memory pressure
5
6 from keras.models import load_model
7 import numpy as np
8 import pandas as pd
9 import matplotlib.pyplot as plt
10 import seaborn as sns
11 from pathlib import Path
12 import joblib
13 from keras.models import load_model as keras_load_model
14 import random
15 from typing import Iterable, Dict, List, Optional, Tuple
```


In [85]:

```
1  # --- functions---
2
3  def load_any_model(path):
4      """
5      Load a model given its file path, handling both Keras (.keras/.h5) and joblib (.joblib/.pkl).
6      """
7      path = Path(path)
8      suf = path.suffix.lower()
9      if suf in {".keras", ".h5", ".hdf5"}:
10         return keras_load_model(path)
11     if suf in {".joblib", ".pkl"}:
12         return joblib.load(path)
13     raise ValueError(f"Unrecognized model suffix for {path.name}")
14
15 def load_models(model_paths: dict):
16     """
17     model_paths: dict like {"Base": "...joblib", "LSTM": "...keras", "CNN": "...keras", "CNN-LSTM": "...keras"}
18     returns: dict {name: model_object}
19     """
20     models = {}
21     for name, path in model_paths.items():
22         models[name] = load_any_model(path)
23     return models
24
25 def build_model_paths(dataset=None, seq_len=None, strategy="last", art_dir=None,
26                       include=("Base", "LSTM", "CNN", "CNN-LSTM")):
27     """
28     Build default file paths for your saved models in <ART_DIR>/models/...
29     Base:      base_linear_{dataset.lower()}_seq{seq_len}_{strategy}.joblib
30     LSTM:      lstm_{dataset.lower()}_seq{seq_len}.keras
31     CNN:       cnn_{dataset.lower()}_seq{seq_len}.keras
32     CNN-LSTM:  cnn_lstm_{dataset.lower()}_seq{seq_len}.keras
33     """
34     if dataset is None:
35         dataset = globals().get("DATASET", "FD001")
36     if seq_len is None:
37         seq_len = globals().get("SEQ_LEN", 30)
38     if art_dir is None:
39         art_dir = globals().get("ART_DIR", Path.cwd() / f"{dataset} data & artefacts")
40
41     models_dir = Path(art_dir) / "models"
```

```

42     ds = dataset.lower()
43
44     paths = {}
45     if "Base" in include:
46         paths["Base"] = str(models_dir / f"base_linear_{ds}_seq{seq_len}_{strategy}.joblib")
47     if "LSTM" in include:
48         paths["LSTM"] = str(models_dir / f"lstm_{ds}_seq{seq_len}.keras")
49     if "CNN" in include:
50         paths["CNN"] = str(models_dir / f"cnn_{ds}_seq{seq_len}.keras")
51     if "CNN-LSTM" in include:
52         paths["CNN-LSTM"] = str(models_dir / f"cnn_lstm_{ds}_seq{seq_len}.keras")
53     return paths
54
55
56 def plot_true_vs_pred(y_true, y_pred, model_name="Model", savepath=None):
57     """
58     Scatter plot: True RUL vs Predicted RUL for a single model.
59
60     Args:
61         y_true (array-like): Ground truth RUL values.
62         y_pred (array-like): Predicted RUL values.
63         model_name (str): Name of the model for labeling.
64         savepath (str or None): If given, save the figure to this path.
65     """
66     plt.figure(figsize=(6, 6))
67     plt.scatter(y_true, y_pred, alpha=0.6, edgecolor="k")
68     max_val = max(y_true.max(), y_pred.max())
69     plt.plot([0, max_val], [0, max_val], "r--", lw=2, label="Ideal")
70     plt.xlabel("True RUL")
71     plt.ylabel("Predicted RUL")
72     plt.title(f"True vs Predicted RUL ({model_name})")
73     plt.legend()
74     plt.grid(True)
75     if savepath:
76         plt.savefig(savepath, dpi=300, bbox_inches="tight")
77     plt.show()
78
79
80
81 def plot_residuals(y_true, y_pred, model_name="Model", kind="box", savepath=None):
82     """
83     Plot residuals (Predicted - True) for a single model.

```

```

84
85     Args:
86         y_true (array-like): Ground truth RUL values.
87         y_pred (array-like): Predicted RUL values.
88         model_name (str): Label for the model.
89         kind (str): "box" or "violin" for plot type.
90         savepath (str or None): If given, save the figure.
91     """
92     errors = y_pred - y_true
93     plt.figure(figsize=(6, 4))
94
95     if kind == "violin":
96         sns.violinplot(y=errors)
97     else:
98         sns.boxplot(y=errors)
99
100    plt.axhline(0, color="r", linestyle="--", lw=2, label="Zero Error")
101    plt.ylabel("Residual (Predicted - True)")
102    plt.title(f"Residual Distribution ({model_name})")
103    plt.legend()
104    plt.grid(True, axis="y")
105    if savepath:
106        plt.savefig(savepath, dpi=300, bbox_inches="tight")
107    plt.show()
108
109 def plot_per_engine_bars(y_true, y_pred, unit_ids, model_name="Model", n_samples=30, savepath=None, seed=42):
110     """
111     Bar plot comparing Actual vs Predicted RUL for a random sample of engines.
112
113     Args:
114         y_true (array-like): Ground truth RUL values (aligned with unit_ids).
115         y_pred (array-like): Predicted RUL values (aligned with unit_ids).
116         unit_ids (array-like): Engine/unit identifiers for each sample.
117         model_name (str): Name of the model for labeling.
118         n_samples (int): Number of engines to randomly sample.
119         savepath (str or None): If given, save the figure.
120         seed (int): Random seed for reproducibility.
121     """
122     random.seed(seed)
123     unique_ids = np.unique(unit_ids)
124     chosen_ids = random.sample(list(unique_ids), min(n_samples, len(unique_ids)))
125

```

```

126     # Collect true & pred for chosen engines
127     true_sample, pred_sample, labels = [], [], []
128     for uid in chosen_ids:
129         mask = unit_ids == uid
130         # Last entry corresponds to the test RUL Label
131         true_sample.append(y_true[mask][-1])
132         pred_sample.append(y_pred[mask][-1])
133         labels.append(str(uid))
134
135     x = np.arange(len(chosen_ids))
136     width = 0.35
137
138     plt.figure(figsize=(12, 6))
139     plt.bar(x - width/2, true_sample, width, label="Actual")
140     plt.bar(x + width/2, pred_sample, width, label="Predicted")
141     plt.xticks(x, labels, rotation=45)
142     plt.xlabel("Engine ID (sampled)")
143     plt.ylabel("RUL")
144     plt.title(f"Actual vs Predicted RUL (Sampled Engines) - {model_name}")
145     plt.legend()
146     plt.tight_layout()
147     if savepath:
148         plt.savefig(savepath, dpi=300, bbox_inches="tight")
149     plt.show()
150 def plot_metric_comparison(metrics_list,
151                             dataset_name: str = "FD001",
152                             savepath: Optional[str] = None):
153     """
154     Grouped bar chart of RMSE/MAE across models.
155     Accepts either:
156     - list of dicts: [{"model": "LSTM", "RMSE": 15.2, "MAE": 11.3}, ...]
157     - DataFrame with columns: model/Model, RMSE, MAE
158     """
159     # Build DataFrame
160     df = metrics_list.copy() if isinstance(metrics_list, pd.DataFrame) else pd.DataFrame(metrics_list)
161     if df.empty:
162         print("No metrics to plot.")
163         return
164
165     # Normalise column names
166     if "model" not in df.columns and "Model" in df.columns:
167         df = df.rename(columns={"Model": "model"})

```

```

168
169     # Validate required columns
170     required = {"model", "RMSE", "MAE"}
171     missing = required - set(df.columns)
172     if missing:
173         raise ValueError(f"Missing columns for plotting: {missing}")
174
175     # Keep only what we need, coerce to numeric
176     df = df[["model", "RMSE", "MAE"]].copy()
177     df[["RMSE", "MAE"]] = df[["RMSE", "MAE"]].astype(float)
178     df = df.set_index("model")
179
180     # Plot
181     ax = df.plot(kind="bar", figsize=(10, 6))
182     ax.set_title(f"RMSE / MAE Comparison - {dataset_name}")
183     ax.set_ylabel("Error")
184     ax.set_xlabel("")
185     ax.grid(True, axis="y", alpha=0.3)
186     plt.xticks(rotation=0)
187     plt.tight_layout()
188     plt.show()

```

Project Module and base set up

In [86]:

```
1  # Project modules
2  import data_loader as dl
3  import pre_processing as pp
4  import evaluator as ev
5  import base_model as base
6  import lstm_model as lstm
7  import cnn_model as cnn
8  import cnn_lstm_model as cnnlstm
9  import plots
10
11  # ---- Paths ----
12  ROOT = Path.cwd()
13  CMAPS = ROOT / "CMaps" # keep correct folder case
14  # ==== Minimal config you tweak next time ====
15  DATASET = "FD001" # <- change this to FD002/FD003/FD004 Later
16  SEQ_LEN = 30 # sliding window
17  MAX_RUL = 130 # RUL clipping
18  VAL_SPLIT = 0.30 # val split by unit
19
20  # Files derived from DATASET (so you edit one line only)
21  TRAIN_PATH = CMAPS / f"train_{DATASET}.txt"
22  TEST_PATH = CMAPS / f"test_{DATASET}.txt"
23  RUL_PATH = CMAPS / f"RUL_{DATASET}.txt"
24
25  # Artifacts folder for this dataset
26  ART_DIR = ROOT / f"{DATASET} data & artefacts"
27  ART_DIR.mkdir(exist_ok=True)
28
29  print(f"backend: torch | dataset: {DATASET}")
30  print("Train:", TRAIN_PATH.name, "| Test:", TEST_PATH.name, "| RUL:", RUL_PATH.name)
```

backend: torch | dataset: FD001

Train: train_FD001.txt | Test: test_FD001.txt | RUL: RUL_FD001.txt

Load & Preprocessing Data

In [87]:

```
1  # --- Load FD001 ---
2  train_df = dl.load_raw_data(CMAPS / f"train_{DATASET}.txt")
3  test_df, rul_df = dl.load_test_data(
4      CMAPS / f"test_{DATASET}.txt",
5      CMAPS / f"RUL_{DATASET}.txt"
6  )
7
8  unit_ids = test_df["unit_number"].values
9
10 print("Loaded.")
11 print("  train_df:", train_df.shape, "  test_df:", test_df.shape, "  rul_df:", rul_df.shape)
12 assert train_df.shape[1] == 26 and test_df.shape[1] == 26
13
14 dl.inspect_data(train_df)
15 pp.summarise_engine_lifespans(train_df, dataset_name=DATASET)
16
17 # -----
18 # 1) TRAIN: make targets first (no Leakage)
19 # -----
20 train_rul = pp.calculate_rul(train_df, max_rul=MAX_RUL)
21
22 # -----
23 # 2) Split by unit BEFORE deciding features/scaling
24 # -----
25 train_split, val_split = pp.split_by_unit(train_rul, test_size=0.2, random_state=42)
26
27 # -----
28 # 3) Decide flat sensors using TRAIN ONLY, then drop same cols from val/test
29 # -----
30 before_cols = list(train_split.columns)
31 train_split_clean = pp.drop_flat_sensors(train_split.copy())
32 after_cols = list(train_split_clean.columns)
33 dropped_cols = [c for c in before_cols if c not in after_cols]
34
35 val_split_clean = val_split.drop(columns=[c for c in dropped_cols if c in val_split.columns]).reset_index(drop=True)
36 test_df_clean = test_df.drop(columns=[c for c in dropped_cols if c in test_df.columns]).reset_index(drop=True)
37
38 # -----
39 # 4) TEST: build true RUL from RUL file, then clip like train
40 # -----
41 last_cycles = test_df_clean.groupby("unit_number")["time_in_cycles"].max()
```

```

42 rul_map = dict(zip(sorted(test_df_clean["unit_number"].unique()), rul_df["RUL"].values))
43 test_df_clean = test_df_clean.copy()
44 test_df_clean["RUL"] = test_df_clean.apply(
45     lambda r: (last_cycles.loc[r["unit_number"]] - r["time_in_cycles"]) + rul_map[r["unit_number"]],
46     axis=1
47 )
48 test_df_clean["RUL"] = np.minimum(test_df_clean["RUL"], MAX_RUL)
49
50 # -----
51 # 5) Scale sensors with ONE scaler fit on TRAIN ONLY (FD001 = single condition)
52 # (avoid standardise_per_condition here to prevent re-fitting on val/test)
53 # -----
54 from sklearn.preprocessing import StandardScaler
55
56 sensor_cols = [c for c in train_split_clean.columns if c.startswith("sensor_measurement")]
57 scaler = StandardScaler().fit(train_split_clean[sensor_cols])
58
59 train_scaled = train_split_clean.copy()
60 val_scaled = val_split_clean.copy()
61 test_scaled = test_df_clean.copy()
62
63 train_scaled[sensor_cols] = scaler.transform(train_scaled[sensor_cols])
64 val_scaled[sensor_cols] = scaler.transform(val_scaled[sensor_cols])
65 test_scaled[sensor_cols] = scaler.transform(test_scaled[sensor_cols])
66
67 # -----
68 # 6) Windowing with your helper
69 # -----
70 X_train, y_train = pp.generate_sliding_windows(train_scaled, seq_len=SEQ_LEN)
71 X_val, y_val = pp.generate_sliding_windows(val_scaled, seq_len=SEQ_LEN)
72 X_test, y_test = pp.generate_sliding_windows(test_scaled, seq_len=SEQ_LEN)
73
74 print("After preprocessing (FD001, no new pp funcs):")
75 print("  Train engines :", train_scaled['unit_number'].nunique())
76 print("  Val engines   :", val_scaled['unit_number'].nunique())
77 print("  X_train shape :", X_train.shape, " y_train:", y_train.shape)
78 print("  X_val  shape :", X_val.shape, " y_val  :", y_val.shape)
79 print("  X_test shape :", X_test.shape, " y_test :", y_test.shape)
80 print("  Dropped sensors:", dropped_cols)
81
82 # --- Save to use in model ---
83 out_npz = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"

```

```
84 pp.save_preprocessed_data(X_train, y_train, X_val, y_val, X_test, y_test, filename=str(out_npz))
85
```

Loaded.

train_df: (20631, 26) test_df: (13096, 26) rul_df: (100, 1)

Shape: (20631, 26)

Unique engines: 100

Missing values:

0

Max cycles per engine:

count 100.000000

mean 206.310000

std 46.342749

min 128.000000

25% 177.000000

50% 199.000000

75% 229.250000


max 362.000000

Name: time_in_cycles, dtype: float64


First 5 rows:

	unit_number	time_in_cycles	op_setting_1	op_setting_2	op_setting_3	sensor_measurement_1	sensor_measurement_2	sensor_measurement_3
0	1	1	-0.0007	-0.0004	100.0	518.67	641.82	1589.70
1	1	2	0.0019	-0.0003	100.0	518.67	642.15	1591.82
2	1	3	-0.0043	0.0003	100.0	518.67	642.35	1587.99
3	1	4	0.0007	0.0000	100.0	518.67	642.35	1582.79
4	1	5	-0.0019	-0.0002	100.0	518.67	642.37	1582.85

5 rows × 26 columns



Distribution of Engine Lifespans (FD001)



Dataset: FD001
Mean cycles to failure: 206.31
Standard deviation: 46.34
Minimum: 128
Maximum: 362
After preprocessing (FD001, no new pp funcs):
 Train engines : 80
 Val engines : 20
 X_train shape : (14241, 30, 15) y_train: (14241,)
 X_val shape : (3490, 30, 15) y_val : (3490,)
 X_test shape : (10196, 30, 15) y_test : (10196,)
 Dropped sensors: ['sensor_measurement_1', 'sensor_measurement_5', 'sensor_measurement_10', 'sensor_measurement_16',
 'sensor_measurement_18', 'sensor_measurement_19']
Data saved to C:\Users\mg020649\Documents\15 - Coding\Msc-Project-main\FD001 data & artefacts\fd001_seq30.npz

=====Train Models=====

Base Model

```
In [88]: 1 from base_model import train_linear_model
2
3 # 1) Load cached windows
4 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
5 X_train, y_train, X_val, y_val, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
6
7 # 2) Convert 3D windows -> 2D feature vectors for baseline
8 # (stick to your existing helper; default strategy='last')
9 X_train_feat = pp.make_feature_vectors_from_windows(X_train, strategy='last')
10 X_val_feat = pp.make_feature_vectors_from_windows(X_val, strategy='last')
11
12 print("Feature shapes (base model):")
13 print(" X_train:", X_train_feat.shape, " y_train:", y_train.shape)
14 print(" X_val  :", X_val_feat.shape, " y_val  :", y_val.shape)
15
16 # 3) Train simple Linear Regression
17 base_model = train_linear_model(X_train_feat, y_train)
18
19 # 4) Save the trained model
20 models_dir = ART_DIR / "models"
21 models_dir.mkdir(parents=True, exist_ok=True)
22
23 model_path = models_dir / f"base_linear_{DATASET.lower()}_seq{SEQ_LEN}_last.joblib"
24 joblib.dump(base_model, model_path)
25
26 print("Saved base model to:", model_path)
```

Feature shapes (base model):

X_train: (14241, 15) y_train: (14241,)

X_val : (3490, 15) y_val : (3490,)

Saved base model to: C:\Users\mg020649\Documents\15 - Coding\Msc-Project-main\FD001 data & artefacts\models\base_line_ar_fd001_seq30_last.joblib

CNN Model

In [89]:

```
1  # --- Train / Save: CNN (FD001) ---
2
3  from cnn_model import build_cnn_model, train_cnn_model
4  import json
5
6  # 1) Load cached 3D windows
7  npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
8  X_train, y_train, X_val, y_val, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
9
10 # 2) Build CNN
11 input_shape = (SEQ_LEN, X_train.shape[2])
12 cnn = build_cnn_model(input_shape)
13
14 # 3) Train (only pass what your module expects!)
15 cnn, history = train_cnn_model(cnn, X_train, y_train, X_val, y_val, epochs=30, batch_size=128)
16
17 # 4) Save model + meta
18 models_dir = ART_DIR / "models"
19 models_dir.mkdir(parents=True, exist_ok=True)
20 model_path = models_dir / f"cnn_{DATASET.lower()}_seq{SEQ_LEN}.keras"
21 cnn.save(model_path)
22
23 meta = {
24     "model_type": "cnn",
25     "dataset": DATASET,
26     "seq_len": int(SEQ_LEN),
27     "features": int(X_train.shape[2]),
28     "epochs": 30,
29     "batch_size": 128
30 }
31 with open(models_dir / f"cnn_{DATASET.lower()}_seq{SEQ_LEN}.meta.json", "w") as f:
32     json.dump(meta, f, indent=2)
33
34 print("Saved CNN model to:", model_path)
```



```

Epoch 1/30
112/112 [=====] - 1s 9ms/step - loss: 2430.6372 - val_loss: 545.6964
Epoch 2/30
112/112 [=====] - 1s 10ms/step - loss: 746.0840 - val_loss: 514.2057
Epoch 3/30
112/112 [=====] - 1s 9ms/step - loss: 680.1778 - val_loss: 475.8336
Epoch 4/30
112/112 [=====] - 1s 10ms/step - loss: 587.3165 - val_loss: 406.2617
Epoch 5/30
112/112 [=====] - 1s 9ms/step - loss: 529.3652 - val_loss: 392.8668
Epoch 6/30
112/112 [=====] - 1s 8ms/step - loss: 475.4832 - val_loss: 355.6811
Epoch 7/30
112/112 [=====] - 1s 8ms/step - loss: 416.3690 - val_loss: 341.5238
Epoch 8/30
112/112 [=====] - 1s 8ms/step - loss: 354.3781 - val_loss: 303.2555
Epoch 9/30
112/112 [=====] - 1s 8ms/step - loss: 317.3850 - val_loss: 287.6030
Epoch 10/30
112/112 [=====] - 1s 8ms/step - loss: 296.0379 - val_loss: 314.2828
Epoch 11/30
112/112 [=====] - 1s 9ms/step - loss: 277.7661 - val_loss: 284.0954
Epoch 12/30
112/112 [=====] - 1s 10ms/step - loss: 256.1288 - val_loss: 281.5417
Epoch 13/30
112/112 [=====] - 1s 9ms/step - loss: 248.9749 - val_loss: 297.4206
Epoch 14/30
112/112 [=====] - 1s 9ms/step - loss: 244.6989 - val_loss: 281.9802
Epoch 15/30
112/112 [=====] - 1s 10ms/step - loss: 222.4565 - val_loss: 289.5781
Epoch 16/30
108/112 [=====>..] - ETA: 0s - loss: 217.2393Restoring model weights from the end of the best e
poch: 12.
112/112 [=====] - 1s 9ms/step - loss: 216.7971 - val_loss: 296.6009
Epoch 16: early stopping
Saved CNN model to: C:\Users\mg020649\Documents\15 - Coding\Msc-Project-main\FD001 data & artefacts\models\cnn_fd001_
seq30.keras

```

LSTM MODEL

In [90]:

```
1 # --- LSTM model (train + save) ---
2
3 from lstm_model import build_lstm_model, train_lstm_model
4 import json
5
6 # 1) Load cached 3D windows
7 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
8 X_train, y_train, X_val, y_val, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
9
10 # 2) Build LSTM (input: [seq_len, n_features])
11 input_shape = (SEQ_LEN, X_train.shape[2])
12 lstm = build_lstm_model(input_shape)
13
14 # 3) Train (pass only what your module expects)
15 # If your train_lstm_model returns (model, history), keep both; if it returns only model, handle that too.
16 result = train_lstm_model(lstm, X_train, y_train, X_val, y_val, epochs=50, batch_size=128)
17 if isinstance(result, tuple):
18     lstm, lstm_history = result
19 else:
20     lstm = result
21     lstm_history = None
22
23 # 4) Save model + minimal metadata
24 models_dir = ART_DIR / "models"
25 models_dir.mkdir(parents=True, exist_ok=True)
26
27 model_path = models_dir / f"lstm_{DATASET.lower()}_seq{SEQ_LEN}.keras"
28 lstm.save(model_path)
29
30 meta = {
31     "model_type": "lstm",
32     "dataset": DATASET,
33     "seq_len": int(SEQ_LEN),
34     "features": int(X_train.shape[2]),
35     "epochs": 50,
36     "batch_size": 128
37 }
38 with open(models_dir / f"lstm_{DATASET.lower()}_seq{SEQ_LEN}.meta.json", "w") as f:
39     json.dump(meta, f, indent=2)
40
```

```
41 print("Saved LSTM model to:", model_path)
```

Epoch 1/50
112/112 [=====] - 7s 41ms/step - loss: 7562.5532 - val_loss: 6861.7861
Epoch 2/50
112/112 [=====] - 4s 35ms/step - loss: 6730.3076 - val_loss: 6315.0210
Epoch 3/50
112/112 [=====] - 4s 35ms/step - loss: 6208.6411 - val_loss: 5828.0205
Epoch 4/50
112/112 [=====] - 4s 36ms/step - loss: 5739.6802 - val_loss: 5387.2583
Epoch 5/50
112/112 [=====] - 4s 35ms/step - loss: 5312.0806 - val_loss: 4980.8403
Epoch 6/50
112/112 [=====] - 4s 38ms/step - loss: 4915.9985 - val_loss: 4604.2236
Epoch 7/50
112/112 [=====] - 4s 35ms/step - loss: 4547.2300 - val_loss: 4253.5591
Epoch 8/50
112/112 [=====] - 4s 34ms/step - loss: 4196.9419 - val_loss: 3928.9534
Epoch 9/50
112/112 [=====] - 4s 34ms/step - loss: 3887.3967 - val_loss: 3628.5186
Epoch 10/50
112/112 [=====] - 4s 34ms/step - loss: 3585.0957 - val_loss: 3348.5396
Epoch 11/50
112/112 [=====] - 4s 33ms/step - loss: 3321.7712 - val_loss: 3091.1445
Epoch 12/50
112/112 [=====] - 4s 34ms/step - loss: 3057.2505 - val_loss: 2850.6289
Epoch 13/50
112/112 [=====] - 4s 34ms/step - loss: 2829.6406 - val_loss: 2626.4160
Epoch 14/50
112/112 [=====] - 4s 33ms/step - loss: 2608.8774 - val_loss: 2425.2144
Epoch 15/50
112/112 [=====] - 4s 33ms/step - loss: 2405.4934 - val_loss: 2225.2141
Epoch 16/50
112/112 [=====] - 4s 34ms/step - loss: 2218.0007 - val_loss: 2045.9836
Epoch 17/50
112/112 [=====] - 4s 34ms/step - loss: 2040.0931 - val_loss: 1879.6632
Epoch 18/50
112/112 [=====] - 4s 34ms/step - loss: 1880.7557 - val_loss: 1730.2860
Epoch 19/50
112/112 [=====] - 4s 34ms/step - loss: 1734.7510 - val_loss: 1586.2383
Epoch 20/50
112/112 [=====] - 4s 34ms/step - loss: 1581.2946 - val_loss: 1453.9030
Epoch 21/50

112/112 [=====] - 4s 34ms/step - loss: 1462.3973 - val_loss: 1334.0830
Epoch 22/50
112/112 [=====] - 4s 33ms/step - loss: 1337.4143 - val_loss: 1222.5902
Epoch 23/50
112/112 [=====] - 4s 33ms/step - loss: 1231.3260 - val_loss: 1121.6599
Epoch 24/50
112/112 [=====] - 4s 34ms/step - loss: 1134.0485 - val_loss: 1024.1525
Epoch 25/50
112/112 [=====] - 4s 34ms/step - loss: 1038.3344 - val_loss: 946.5021
Epoch 26/50
112/112 [=====] - 4s 34ms/step - loss: 947.4596 - val_loss: 864.1041
Epoch 27/50
112/112 [=====] - 4s 35ms/step - loss: 876.8146 - val_loss: 788.6476
Epoch 28/50
112/112 [=====] - 4s 35ms/step - loss: 806.3655 - val_loss: 723.6217
Epoch 29/50
112/112 [=====] - 4s 36ms/step - loss: 739.7886 - val_loss: 667.5870
Epoch 30/50
112/112 [=====] - 4s 36ms/step - loss: 675.3588 - val_loss: 626.8652
Epoch 31/50
112/112 [=====] - 4s 34ms/step - loss: 626.2643 - val_loss: 568.7270
Epoch 32/50
112/112 [=====] - 4s 34ms/step - loss: 574.7999 - val_loss: 515.9753
Epoch 33/50
112/112 [=====] - 4s 34ms/step - loss: 531.2156 - val_loss: 498.0333
Epoch 34/50
112/112 [=====] - 4s 35ms/step - loss: 487.1048 - val_loss: 439.2585
Epoch 35/50
112/112 [=====] - 4s 34ms/step - loss: 451.7163 - val_loss: 414.2429
Epoch 36/50
112/112 [=====] - 4s 34ms/step - loss: 411.0243 - val_loss: 389.1003
Epoch 37/50
112/112 [=====] - 4s 34ms/step - loss: 381.2126 - val_loss: 369.4634
Epoch 38/50
112/112 [=====] - 4s 34ms/step - loss: 352.9323 - val_loss: 334.5475
Epoch 39/50
112/112 [=====] - 4s 35ms/step - loss: 336.7275 - val_loss: 324.1405
Epoch 40/50
112/112 [=====] - 4s 35ms/step - loss: 304.2496 - val_loss: 319.7469
Epoch 41/50
112/112 [=====] - 4s 35ms/step - loss: 283.5230 - val_loss: 304.2993
Epoch 42/50

```
112/112 [=====] - 4s 34ms/step - loss: 259.3516 - val_loss: 291.6107
Epoch 43/50
112/112 [=====] - 4s 35ms/step - loss: 240.0966 - val_loss: 289.1374
Epoch 44/50
112/112 [=====] - 4s 35ms/step - loss: 229.1109 - val_loss: 270.8332
Epoch 45/50
112/112 [=====] - 4s 34ms/step - loss: 210.8021 - val_loss: 261.8469
Epoch 46/50
112/112 [=====] - 4s 35ms/step - loss: 201.1647 - val_loss: 248.1543
Epoch 47/50
112/112 [=====] - 4s 34ms/step - loss: 190.7978 - val_loss: 270.1429
Epoch 48/50
112/112 [=====] - 4s 34ms/step - loss: 183.3460 - val_loss: 286.1623
Epoch 49/50
112/112 [=====] - 4s 34ms/step - loss: 168.6442 - val_loss: 265.0956
Saved LSTM model to: C:\Users\mg020649\Documents\15 - Coding\Msc-Project-main\FD001 data & artefacts\models\lstm_fd00
1_seq30.keras
```

CNN_LSTM Model

In [91]:

```
1 # --- CNN-LSTM model (train + save) ---
2
3 from cnn_lstm_model import build_cnn_lstm_model, train_cnn_lstm_model
4 import json
5
6 # 1) Load cached 3D windows
7 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
8 X_train, y_train, X_val, y_val, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
9
10 # 2) Build model (input: [seq_len, n_features])
11 input_shape = (SEQ_LEN, X_train.shape[2])
12 cnnlstm = build_cnn_lstm_model(input_shape)
13
14 # 3) Train (handle: returns History, returns Model, or returns (Model, History))
15 train_ret = train_cnn_lstm_model(
16     cnnlstm, X_train, y_train, X_val, y_val,
17     epochs=50, batch_size=128
18 )
19
20 # Normalize outputs
21 cnnlstm_history = None
22 model_to_save = cnnlstm # fallback: the model we built
23
24 # Case A: (model, history)
25 if isinstance(train_ret, tuple) and len(train_ret) >= 1:
26     if hasattr(train_ret[0], "save"):
27         model_to_save = train_ret[0]
28     if len(train_ret) >= 2 and hasattr(train_ret[1], "history"):
29         cnnlstm_history = train_ret[1]
30
31 # Case B: just a model
32 elif hasattr(train_ret, "save"):
33     model_to_save = train_ret
34
35 # Case C: just a History
36 elif hasattr(train_ret, "history"):
37     cnnlstm_history = train_ret
38
39 # 4) Save model + minimal metadata
40 models_dir = ART_DIR / "models"
41 models_dir.mkdir(parents=True, exist_ok=True)
```

```
42
43 model_path = models_dir / f"cnn_lstm_{DATASET.lower()}_seq{SEQ_LEN}.keras"
44 model_to_save.save(model_path)
45
46 meta = {
47     "model_type": "cnn_lstm",
48     "dataset": DATASET,
49     "seq_len": int(SEQ_LEN),
50     "features": int(X_train.shape[2]),
51     "epochs": 50,           # keep in sync with your fit(...)
52     "batch_size": 128       # keep in sync with your fit(...)
53 }
54 with open(models_dir / f"cnn_lstm_{DATASET.lower()}_seq{SEQ_LEN}.meta.json", "w") as f:
55     json.dump(meta, f, indent=2)
56
57 # (Optional) persist training history if available
58 if cnnlstm_history:
59     hist_path = models_dir / f"cnn_lstm_{DATASET.lower()}_seq{SEQ_LEN}.history.json"
60     with open(hist_path, "w") as f:
61         json.dump(cnnlstm_history.history, f, indent=2)
62
63
64 print("Saved CNN-LSTM model to:", model_path)
```

Epoch 1/50
112/112 [=====] - 6s 32ms/step - loss: 5221.5181 - mae: 60.3353 - val_loss: 2435.6050 - val_mae: 40.1090
Epoch 2/50
112/112 [=====] - 3s 29ms/step - loss: 1291.7435 - mae: 28.2668 - val_loss: 657.3648 - val_mae: 22.8221
Epoch 3/50
112/112 [=====] - 3s 29ms/step - loss: 599.5092 - mae: 18.9909 - val_loss: 405.6673 - val_mae: 17.2569
Epoch 4/50
112/112 [=====] - 3s 30ms/step - loss: 452.6629 - mae: 16.1052 - val_loss: 274.5971 - val_mae: 13.1704
Epoch 5/50
112/112 [=====] - 3s 29ms/step - loss: 376.0709 - mae: 14.6896 - val_loss: 293.0319 - val_mae: 12.5061
Epoch 6/50
112/112 [=====] - 3s 29ms/step - loss: 319.1649 - mae: 13.4552 - val_loss: 299.3013 - val_mae: 13.9737
Epoch 7/50
112/112 [=====] - 3s 29ms/step - loss: 298.9593 - mae: 13.0250 - val_loss: 328.6744 - val_mae: 14.2838
Epoch 8/50
112/112 [=====] - 3s 29ms/step - loss: 280.8509 - mae: 12.6068 - val_loss: 320.2022 - val_mae: 13.7744
Saved CNN-LSTM model to: C:\Users\mg020649\Documents\15 - Coding\Msc-Project-main\FD001 data & artefacts\models\cnn_lstm_fd001_seq30.keras

===== End of Training =====

===== Model Analysis =====

LOAD MODEL

In [92]:

```
1 # LOAD baseline MODEL
2 model_paths = build_model_paths() # will only load what exists
3 print("Will load:", model_paths)
4
5 models = load_models(model_paths)
6 print("Loaded models:", list(models.keys()))
```

Will load: {'Base': 'C:\\Users\\mg020649\\Documents\\15 - Coding\\Msc-Project-main\\FD001 data & artefacts\\models\\base_linear_fd001_seq30_last.joblib', 'LSTM': 'C:\\Users\\mg020649\\Documents\\15 - Coding\\Msc-Project-main\\FD001 data & artefacts\\models\\lstm_fd001_seq30.keras', 'CNN': 'C:\\Users\\mg020649\\Documents\\15 - Coding\\Msc-Project-main\\FD001 data & artefacts\\models\\cnn_fd001_seq30.keras', 'CNN-LSTM': 'C:\\Users\\mg020649\\Documents\\15 - Coding\\Msc-Project-main\\FD001 data & artefacts\\models\\cnn_lstm_fd001_seq30.keras'}

Loaded models: ['Base', 'LSTM', 'CNN', 'CNN-LSTM']


=====ALL MODEL LOADED=====

Base build Analysis


In [93]:

```
1  # plot_true_vs_pred
2
3  npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
4
5  _, _, _, _, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
6
7  X_test_base = pp.make_feature_vectors_from_windows(X_test, strategy="last")
8
9  y_pred_base = models["Base"].predict(X_test_base).ravel()
10
11  y_true = y_test
12
13  plot_true_vs_pred(y_true, y_pred_base, model_name="Base")
```

True vs Predicted RUL (Base)




```
In [94]: 1 # Now plot residuals  
2 plot_residuals(y_true, y_pred_base, model_name="Base", kind="violin")
```



In [95]:

```
1 # Build unit IDs aligned to the windowed test data (X_test, y_test)
2 ids = []
3 for uid, grp in test_scaled.groupby("unit_number"):
4     grp = grp.sort_values("time_in_cycles")
5     n = len(grp) - SEQ_LEN + 1
6     if n > 0:
7         ids.extend([uid] * n)
8 unit_ids_aligned = np.array(ids)
9
10 # Now plot
11 plot_per_engine_bars(y_true, y_pred_base, unit_ids_aligned, model_name="Base")
12
```



===== End of Base TEST =====

CNN Model Analysis

In [96]:


```
1 # CNN Prediction
2 from cnn_model import predict_cnn_model
3 y_pred_cnn = predict_cnn_model(models["CNN"], X_test)
```

319/319 [=====] - 1s 2ms/step


In [97]:

```
1 # plot_true_vs_pred
2
3 # can be rid, brought as don't want to run model build code again
4 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
5 _, _, _, _, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
6
7 y_true = y_test
8
9 plot_true_vs_pred(y_true, y_pred_cnn, model_name="CNN")
```

True vs Predicted RUL (CNN)



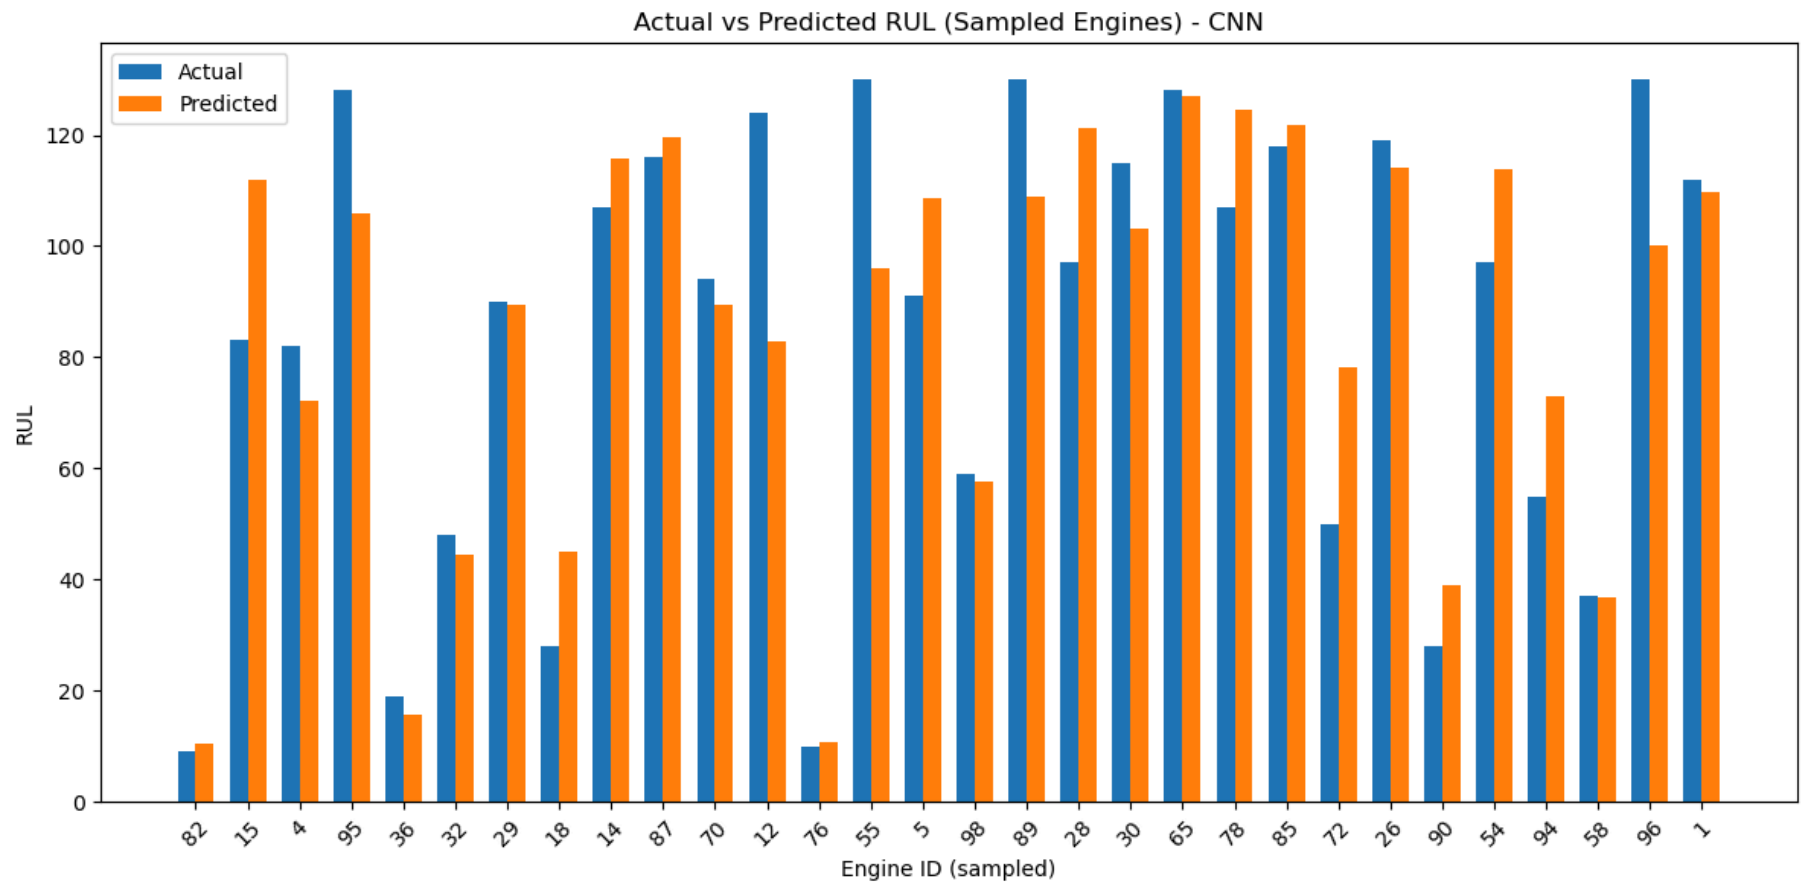
```
In [98]: 1 plot_residuals(y_true, y_pred_cnn, model_name="CNN", kind="violin")
```



```

In [99]: 1 # Build unit IDs aligned to the windowed test data (X_test, y_test)
2 ids = []
3 for uid, grp in test_scaled.groupby("unit_number"):
4     grp = grp.sort_values("time_in_cycles")
5     n = len(grp) - SEQ_LEN + 1
6     if n > 0:
7         ids.extend([uid] * n)
8 unit_ids_aligned = np.array(ids)
9
10 # Now plot
11 plot_per_engine_bars(y_true, y_pred_cnn, unit_ids_aligned, model_name="CNN")

```



===== End of CNN TEST =====

LSTM ANALYSIS

In [100]:

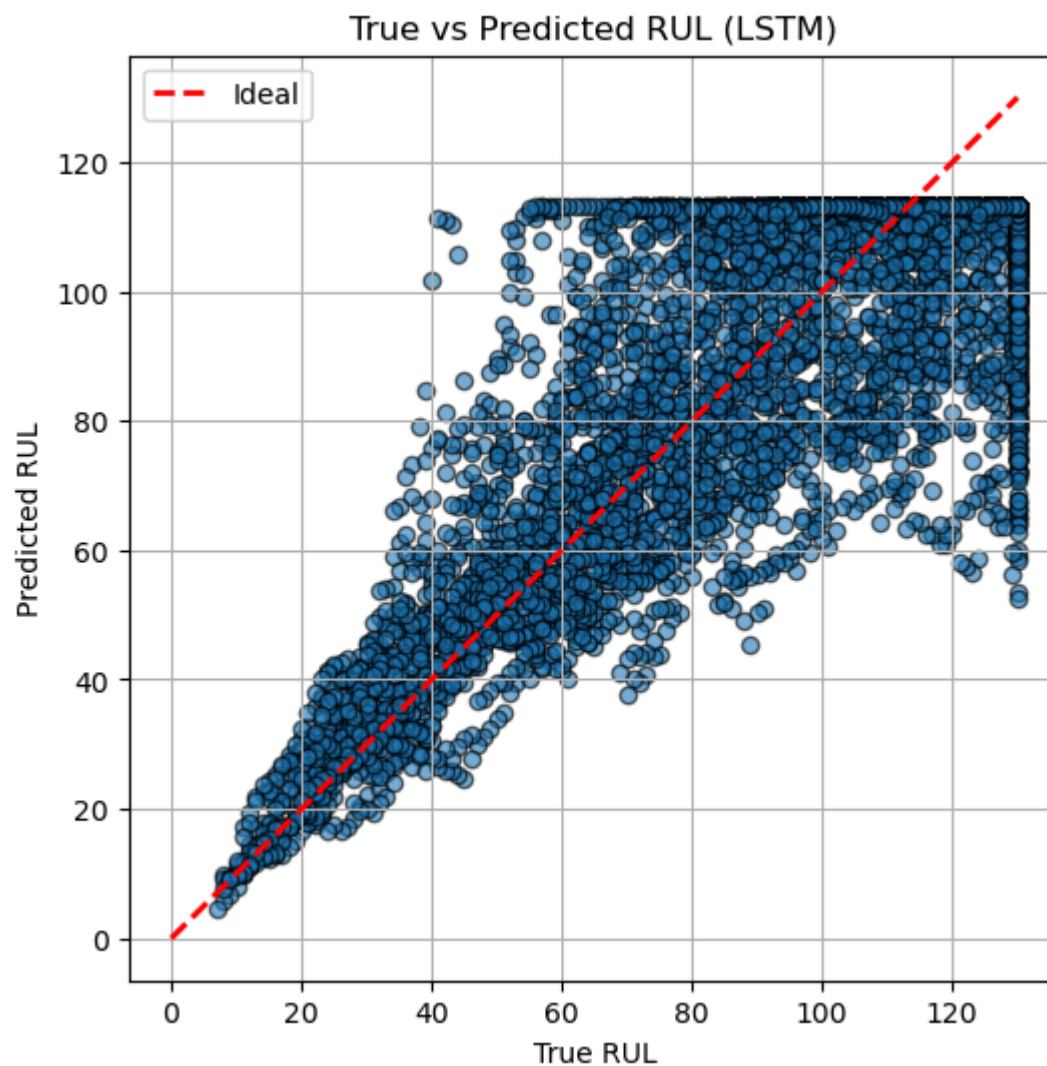
```
1 # CNN Prediction
2 from lstm_model import predict_lstm_model
3 y_pred_lstm = predict_lstm_model(models["LSTM"], X_test)
```

319/319 [=====] - 2s 6ms/step


```

In [101]: 1 # plot_true_vs_pred
          2
          3 # can be rid, brought as don't want to run model build code again
          4 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
          5 _, _, _, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
          6
          7 y_true = y_test
          8 plot_true_vs_pred(y_true, y_pred_lstm, model_name="LSTM")

```

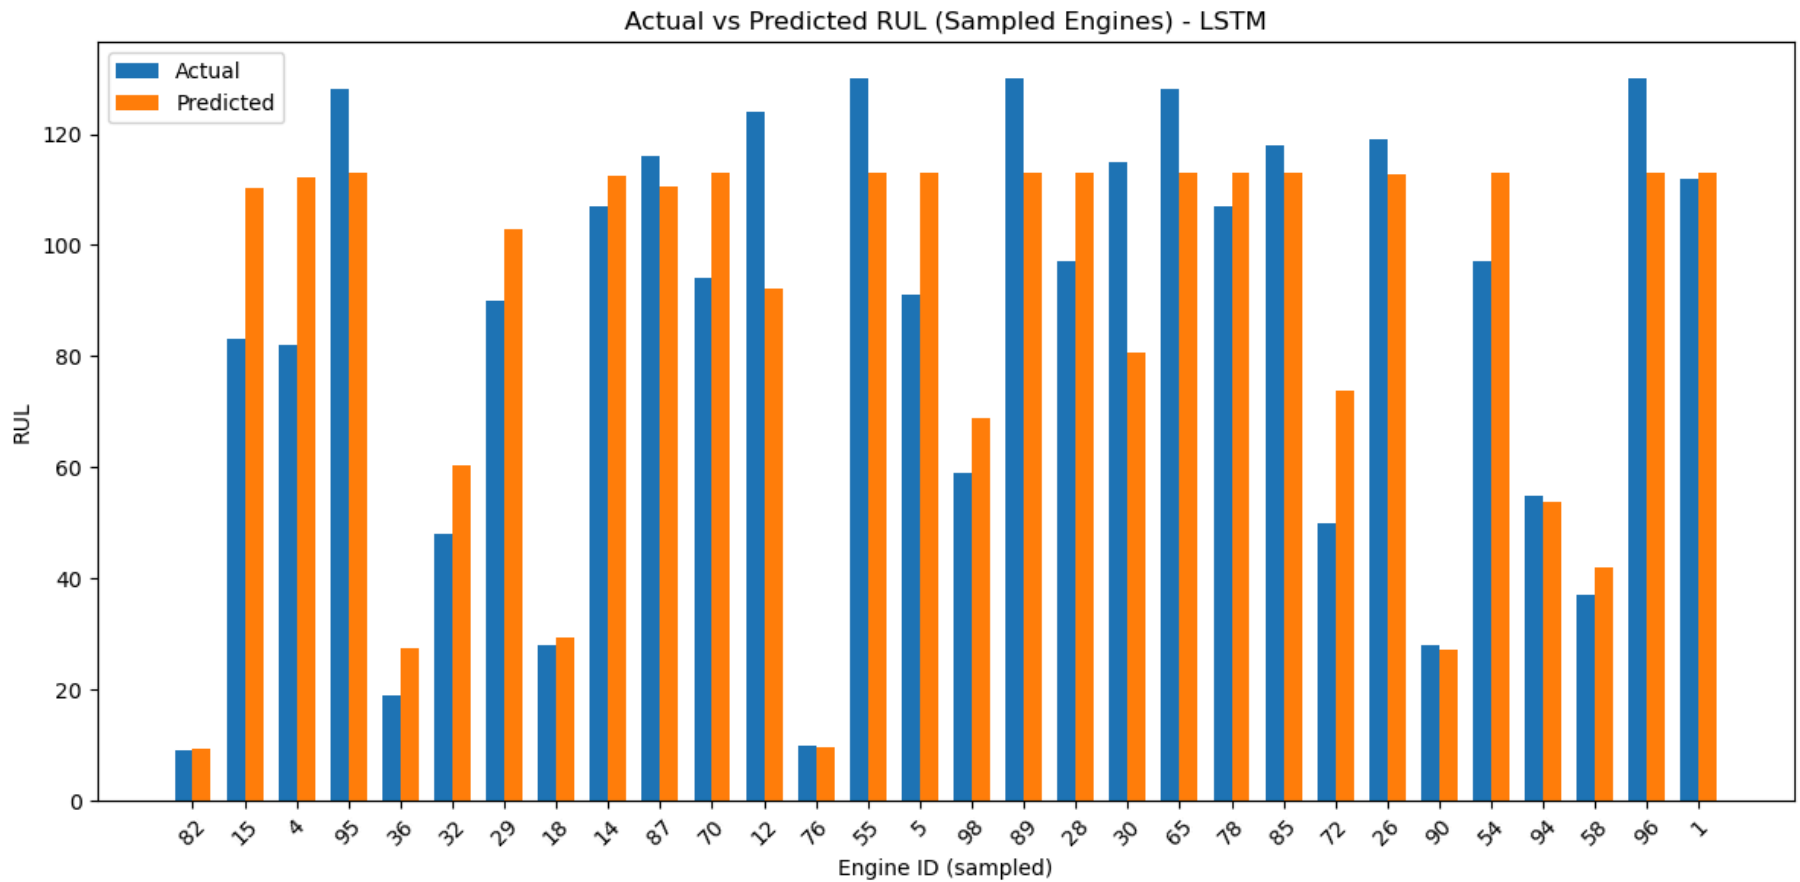


```
In [102]: 1 plot_residuals(y_true, y_pred_lstm, model_name="LSTM", kind="violin")
```



In [103]:

```
1 # Build unit IDs aligned to the windowed test data (X_test, y_test)
2 ids = []
3 for uid, grp in test_scaled.groupby("unit_number"):
4     grp = grp.sort_values("time_in_cycles")
5     n = len(grp) - SEQ_LEN + 1
6     if n > 0:
7         ids.extend([uid] * n)
8 unit_ids_aligned = np.array(ids)
9
10 # Now plot
11 plot_per_engine_bars(y_true, y_pred_lstm, unit_ids_aligned, model_name="LSTM")
12
```



===== End of LSTM TEST =====


LSTM_CNN Build Analysis

In [104]:


```
1 # CNN Prediction
2 from cnn_lstm_model import predict_cnn_lstm_model
3 y_pred_cnn_lstm = predict_cnn_lstm_model(models["CNN-LSTM"], X_test)
```

In [105]:

```
1 # plot_true_vs_pred
2
3 # can be rid, brought as don't want to run model build code again
4 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
5 _, _, _, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
6
7 y_true = y_test
8 plot_true_vs_pred(y_true, y_pred_cnn_lstm, model_name="CNN-LSTM")
```



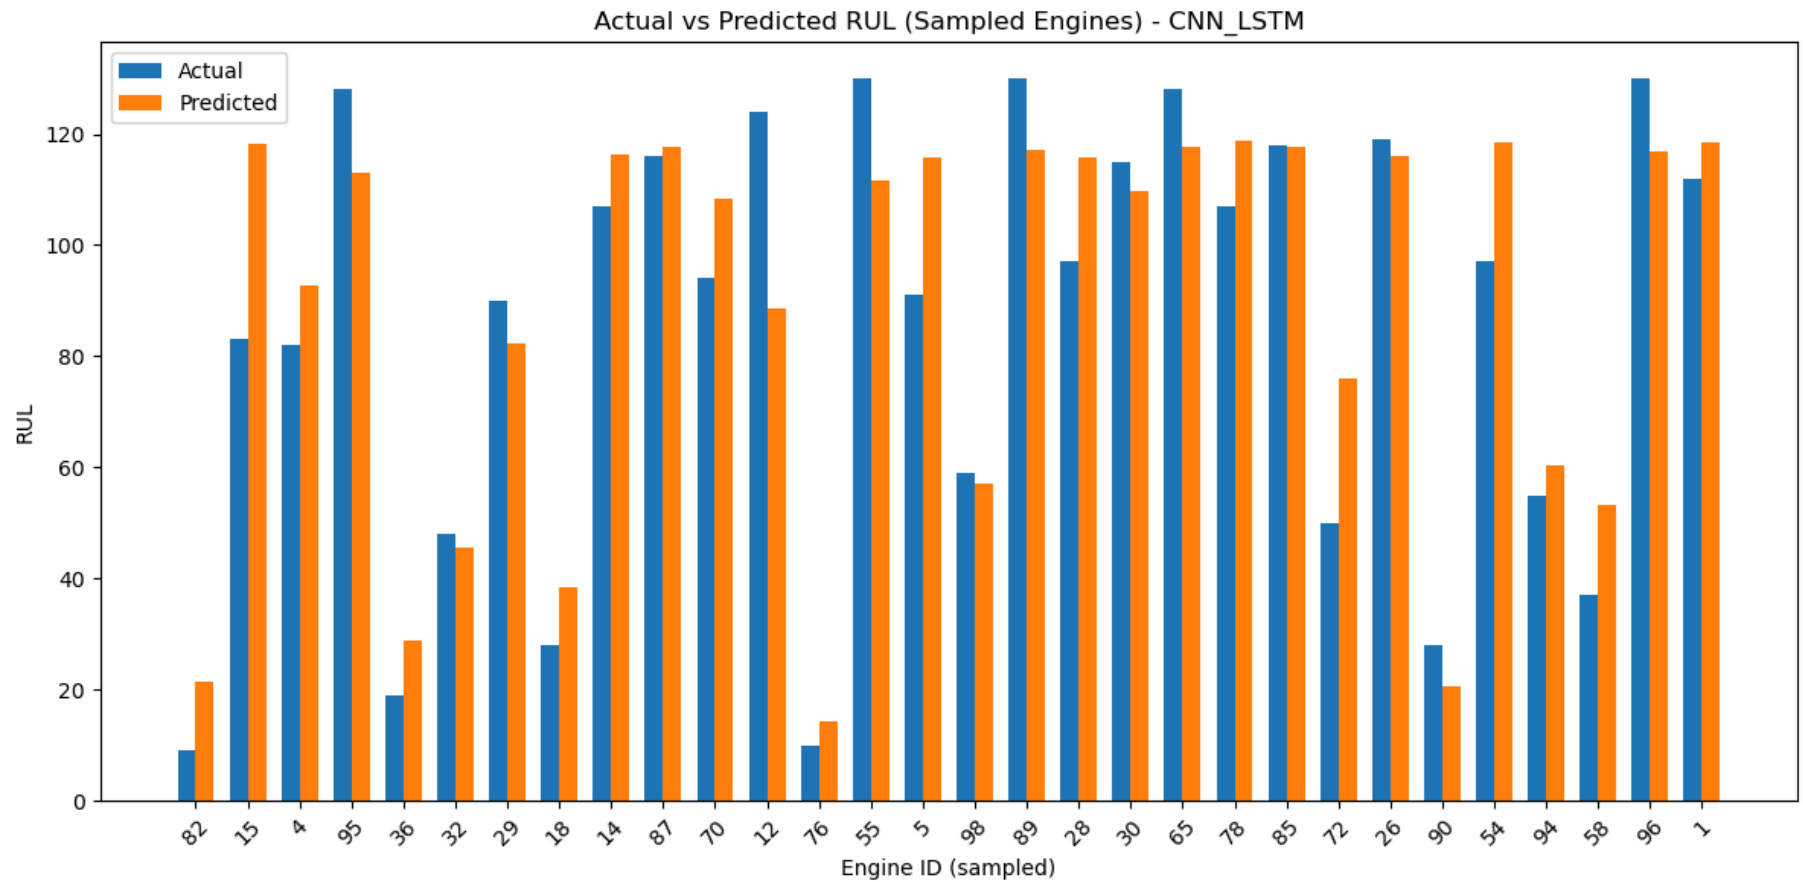
```
In [106]: 1 plot_residuals(y_true, y_pred_cnn_lstm, model_name="CNN-LSTM", kind="violin")
```



```

In [107]: 1 # Build unit IDs aligned to the windowed test data (X_test, y_test)
2 ids = []
3 for uid, grp in test_scaled.groupby("unit_number"):
4     grp = grp.sort_values("time_in_cycles")
5     n = len(grp) - SEQ_LEN + 1
6     if n > 0:
7         ids.extend([uid] * n)
8 unit_ids_aligned = np.array(ids)
9
10 # Now plot
11 plot_per_engine_bars(y_true, y_pred_cnn_lstm, unit_ids_aligned, model_name="CNN_LSTM")
12

```



===== End of CNN_LSTM TEST =====

===== Evaluation Section =====

===== RMSE & MAE =====

base model

```
In [108]: 1 from evaluator import evaluate_model
          2
          3 # Load preprocessed test set
          4 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
          5 _, _, _, _, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
          6
          7 # convert if using base model (needs 2D vectors)
          8 X_test_feat = pp.make_feature_vectors_from_windows(X_test, strategy="last")
          9
         10 # Load trained base model
         11 model_path = ART_DIR / "models" / f"base_linear_{DATASET.lower()}_seq{SEQ_LEN}_last.joblib"
         12 base_model = joblib.load(model_path)
         13
         14 # predict & evaluate
         15 y_base_pred = base_model.predict(X_test_feat)
         16 evaluate_model(y_test, y_base_pred, model_name="Base Linear Model")
```

Base Linear Model Evaluation:

RMSE: 23.2810

MAE : 18.8226

```
Out[108]: {'model': 'Base Linear Model', 'RMSE': 23.280973, 'MAE': 18.8226}
```

CNN model

```
In [109]: 1 # --- Predict & evaluate: CNN on FD001 test ---
2
3 from evaluator import evaluate_model
4
5 # 1) Load cached test windows
6 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
7 _, _, _, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
8
9 # 2) Load saved CNN
10 cnn_path = ART_DIR / "models" / f"cnn_{DATASET.lower()}_seq{SEQ_LEN}.keras"
11 cnn_model = load_model(cnn_path)
12
13 # 3) Predict (CNN expects 3D windows)
14 y_cnn_pred = cnn_model.predict(X_test, verbose=0).squeeze()
15
16 # 4) Evaluate
17 evaluate_model(y_test, y_cnn_pred, model_name="CNN")
```

CNN Evaluation:

RMSE: 19.5039

MAE : 15.2885

```
Out[109]: {'model': 'CNN', 'RMSE': 19.503937, 'MAE': 15.288462}
```

LSTM model

```
In [110]: 1 # --- Predict & evaluate: LSTM on FD001 test ---
2
3 from evaluator import evaluate_model
4
5 # 1) Load cached test windows
6 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
7 _, _, _, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
8
9 # 2) Load saved LSTM
10 lstm_path = ART_DIR / "models" / f"lstm_{DATASET.lower()}_seq{SEQ_LEN}.keras"
11 lstm_model = load_model(lstm_path)
12
13 # 3) Predict (LSTM expects 3D input)
14 y_lstm_pred = lstm_model.predict(X_test, verbose=0).squeeze()
15
16 # 4) Evaluate
17 evaluate_model(y_test, y_lstm_pred, model_name="LSTM")
```

LSTM Evaluation:

RMSE: 18.6124

MAE : 15.7869

```
Out[110]: {'model': 'LSTM', 'RMSE': 18.612432, 'MAE': 15.786917}
```

CNN_LSTM model


```
In [111]: 1 # --- Predict & evaluate: LSTM on FD001 test ---
2
3 from evaluator import evaluate_model
4
5 # 1) Load cached test windows
6 npz_path = ART_DIR / f"{DATASET.lower()}_seq{SEQ_LEN}.npz"
7 _, _, _, X_test, y_test = pp.load_preprocessed_data(str(npz_path))
8
9 # 2) Load saved LSTM
10 lstm_path = ART_DIR / "models" / f"cnn_lstm_{DATASET.lower()}_seq{SEQ_LEN}.keras"
11 lstm_model = load_model(lstm_path)
12
13 # 3) Predict (LSTM expects 3D input)
14 y_cnn_lstm_pred = lstm_model.predict(X_test, verbose=0).squeeze()
15
16 # 4) Evaluate
17 evaluate_model(y_test, y_cnn_lstm_pred, model_name="CNN_LSTM")
```

CNN_LSTM Evaluation:

RMSE: 19.1321

MAE : 15.2806

```
Out[111]: {'model': 'CNN_LSTM', 'RMSE': 19.132107, 'MAE': 15.280622}
```

```
In [112]: 1 from evaluator import evaluate_model
2
3 # build tables
4
5 res_cnn      = evaluate_model(y_test, y_pred_cnn,      model_name="CNN")
6 res_lstm     = evaluate_model(y_test, y_lstm_pred,     model_name="LSTM")
7 res_cnnlstm  = evaluate_model(y_test, y_cnn_lstm_pred, model_name="CNN-LSTM")
8 res_base     = evaluate_model(y_test, y_base_pred,     model_name="Base")
9
10 metrics = [res_cnn, res_lstm, res_cnnlstm, res_base]
11 metrics
```

CNN Evaluation:

RMSE: 19.5039

MAE : 15.2885

LSTM Evaluation:

RMSE: 18.6124

MAE : 15.7869

CNN-LSTM Evaluation:

RMSE: 19.1321

MAE : 15.2806

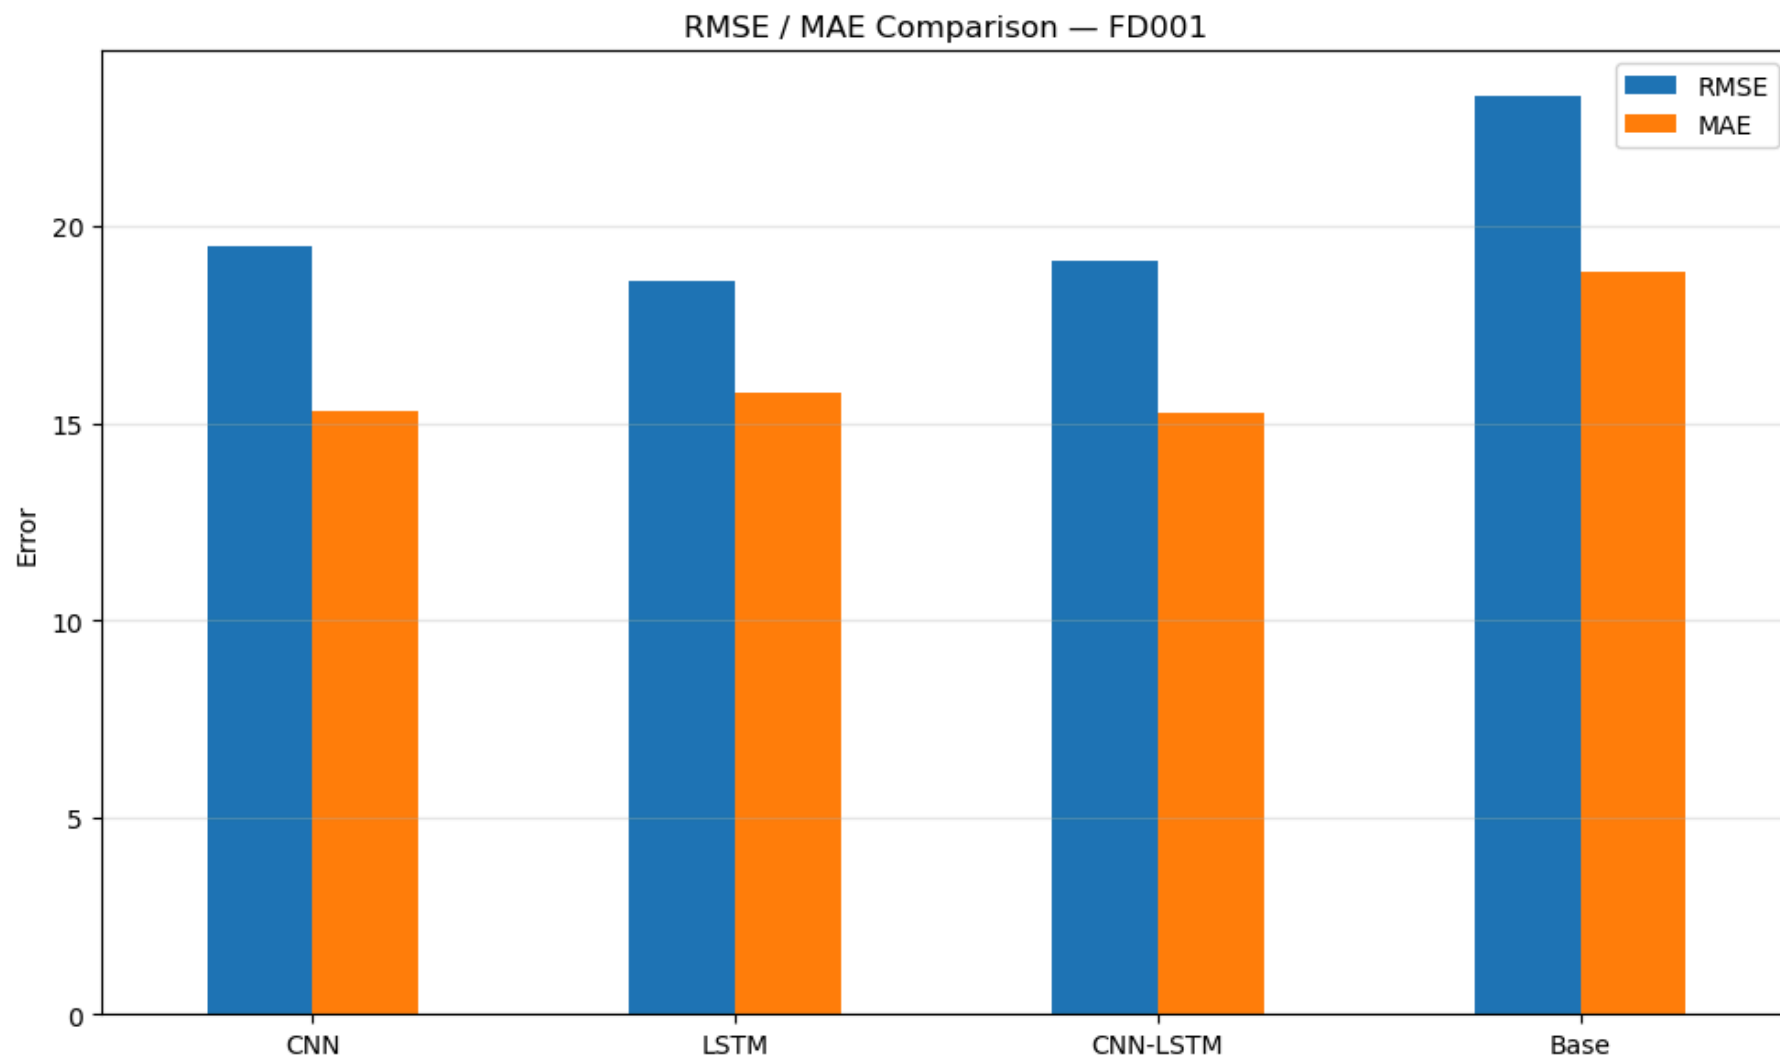
Base Evaluation:

RMSE: 23.2810

MAE : 18.8226

```
Out[112]: [{'model': 'CNN', 'RMSE': 19.503937, 'MAE': 15.288462},
{'model': 'LSTM', 'RMSE': 18.612432, 'MAE': 15.786917},
{'model': 'CNN-LSTM', 'RMSE': 19.132107, 'MAE': 15.280622},
{'model': 'Base', 'RMSE': 23.280973, 'MAE': 18.8226}]
```

```
In [113]: 1 plot_metric_comparison(metrics, dataset_name=DATASET)
```



```
In [114]: 1 metrics_df = pd.DataFrame([res_cnn, res_lstm, res_cnnlstm, res_base])[
2         ["model", "RMSE", "MAE"]
3         ].rename(columns={"model": "Model"})
4
5 metrics_df = metrics_df.sort_values("RMSE").reset_index(drop=True).round(2)
6 display(metrics_df)
```

	Model	RMSE	MAE
0	LSTM	18.610001	15.79
1	CNN-LSTM	19.129999	15.28
2	CNN	19.500000	15.29
3	Base	23.280001	18.82

```
In [115]: 1 metrics_df.to_csv(ART_DIR / f"{DATASET.lower()}_metrics_seq{SEQ_LEN}.csv", index=False)
```